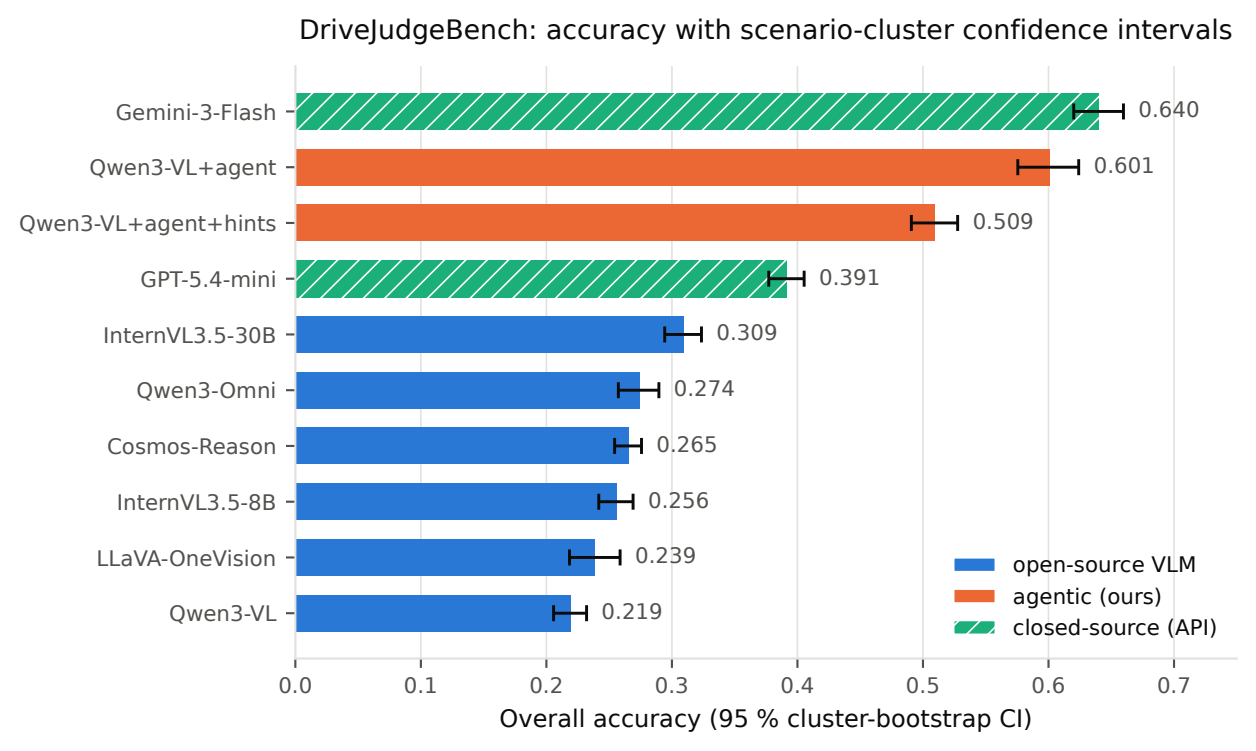
DriveJudgeBench: accuracy with scenario-cluster confidence intervals
Gemini-3-Flash
Qwen3-VL+agent
Qwen3-VL+agent+hints
GPT-5.4-mini
InternVL3.5-30B
Qwen3-Omni
Cosmos-Reason
InternVL3.5-8B
LLaVA-OneVision
Qwen3-VL
0.640
0.601
0.509
0.391
0.309
0.274
0.265
0.256
0.239
0.219
open-source VLM
agentic (ours)
closed-source (API)
0.0
0.1
0.2
0.3
0.4
0.5
0.6
0.7
Overall accuracy (95 % cluster-bootstrap CI)